

Unsupervised learning: Clustering

Pattern Recognition and Machine Learning - 2017

Davide Abati

September 28th, 2017

University of Modena and Reggio Emilia

K-means

Spectral clustering

K-means

- Is a partitional clustering model
 - splits data $\{x_i\}_1^n$ into k disjoint sets
 - the number of sets k has to be provided as input
- solves the following optimization problem:

$$\arg \min_{\{c_1, \dots, c_k\}} = \sum_{j=1}^k \sum_{i=1}^n \mathbf{I}(i, j) \|x_i - c_j\|^2$$

$$\mathbf{I}(i, j) = \begin{cases} 1, & x_i \text{ belongs to cluster } j \\ 0, & \text{otherwise} \end{cases}$$

The problem is NP-hard. A simple heuristic algorithm can be employed to converge to a *local* minimum:

- Initialize k centers randomly
- Repeat until convergence:
 - assign each example to the closest center (i.e. lower euclidean distance)
 - re-estimate centers as the mean of their clusters

Try to implement it from scratch!

K-means can be employed for image segmentation, simply by grouping pixels in the color space. You can also add coordinates to each pixel to obtain a smooth output.

Image




Segmentation




Try it!

- it can get stuck into bad local minima
 - OPTIONAL: run the algorithm many times and choose the most recurrent solution
- can only be employed in spaces where the mean operation is defined
- due to its cost function, it can only cope with compact ball-shaped clusters

GT



Result



Spectral clustering

Clustering model based on the spectral graph theory.

- build a graph over examples, representing it with the adjacency matrix A


$$A_{i,j} = e^{-\frac{\sum_{k=1}^d ||x_i^k - x_j^k||^2}{\sigma^2}}$$

- build the degree matrix D of the graph. It is a diagonal matrix holding for each element the sum of the incoming edges.
- compute the normalized laplacian L


$$L = I - D^{-0.5} A D^{-0.5} \quad (1)$$

- Compute the eigenvectors and sort them for increasing eigenvalues
- Choose the eigenvectors from the second to the desired number of clusters
- Those eigenvectors provide a representation of data in a fancy embedding space: run K-means over such eigenvectors.




GT



Result



Beware: σ counts (in large amounts!)



- `datasets.gaussian_dataset`

```
data, cl = gaussians_dataset(3, [100,100,70], [[1, 1],[-4, 6],[8, 8]], [[1, 1],[3, 3],[1, 1]])
```

- `datasets.two_moon_dataset`

```
data, cl = two_moon_dataset(n_samples=300, noise=0.1)
```

- `matplotlib.pyplot.plot`
- `matplotlib.pyplot.scatter`
- `scipy.linalg.fractional_matrix_power`
- `numpy.linalg.eig`
- `numpy.argsort`